

Problem Statement:

A real estate agent want to help to predict the house price for regions in USA.He gave us the dataset to work on to use Linear Regression modelCreate a Model that helps him to estimate of what the house would sell for

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df=pd.read_csv("13_placement.csv")
df
```

```
Out[2]:
```

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0
...
995	8.87	44.0	1
996	9.12	65.0	1
997	4.89	34.0	0
998	8.62	46.0	1
999	4.90	10.0	1

1000 rows × 3 columns

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   cgpa                   1000 non-null   float64
1   placement_exam_marks  1000 non-null   float64
2   placed                 1000 non-null   int64
dtypes: float64(2), int64(1)
memory usage: 23.6 KB
```

```
In [4]: df.head()
```

Out[4]:

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0

Data cleaning and Pre-Processing

In [5]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   cgpa                  1000 non-null  float64
1   placement_exam_marks 1000 non-null  float64
2   placed                1000 non-null  int64
dtypes: float64(2), int64(1)
memory usage: 23.6 KB
```

In [6]:

```
df.describe()
```

Out[6]:

	cgpa	placement_exam_marks	placed
count	1000.000000	1000.000000	1000.000000
mean	6.961240	32.225000	0.489000
std	0.615898	19.130822	0.500129
min	4.890000	0.000000	0.000000
25%	6.550000	17.000000	0.000000
50%	6.960000	28.000000	0.000000
75%	7.370000	44.000000	1.000000
max	9.120000	100.000000	1.000000

In [7]:

```
df.dropna(axis='columns')
```

Out[7]:

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1

	cgpa	placement_exam_marks	placed
4	7.23	17.0	0
...
995	8.87	44.0	1
996	9.12	65.0	1
997	4.89	34.0	0
998	8.62	46.0	1
999	4.90	10.0	1

1000 rows × 3 columns

```
In [8]: a = df.dropna(axis='columns')
a.columns
```

```
Out[8]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')
```

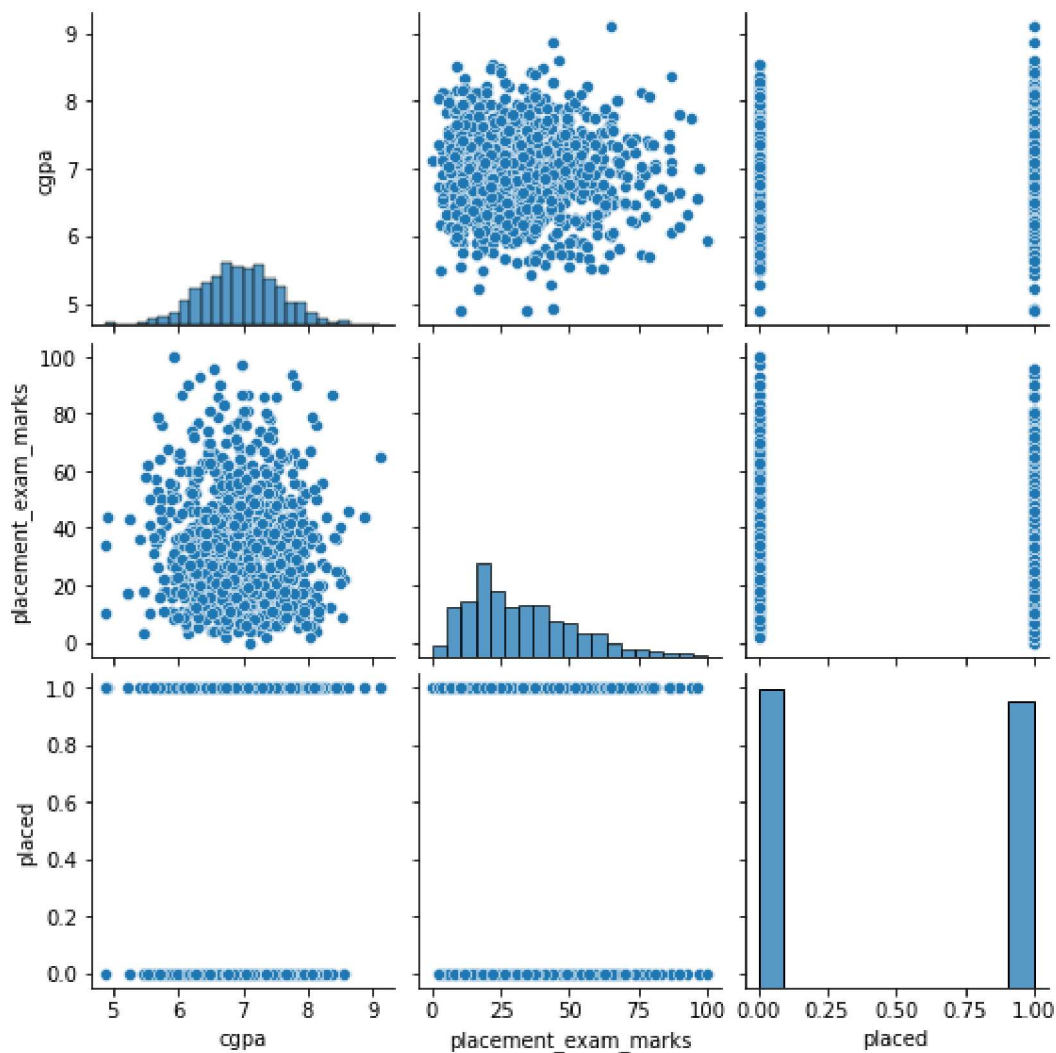
```
In [9]: df.columns
```

```
Out[9]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')
```

EDA and VISUALIZATION

```
In [10]: sns.pairplot(df)
```

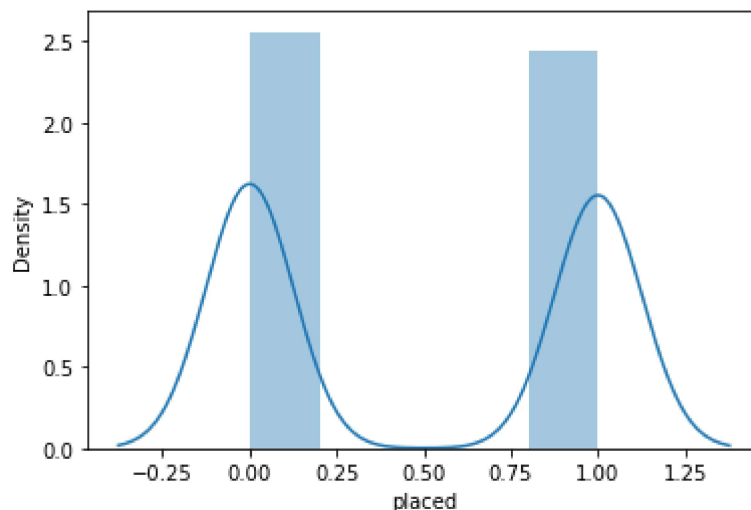
```
Out[10]: <seaborn.axisgrid.PairGrid at 0x1a382264550>
```



```
In [12]: sns.distplot(df['placed'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[12]: <AxesSubplot:xlabel='placed', ylabel='Density'>
```




```
In [13]: df1=df[['cgpa', 'placement_exam_marks', 'placed']]
```

Plot Using Heat Map

```
In [14]: sns.heatmap(df1.corr())
```

Out[14]: <AxesSubplot:>



To Train The Model-Model Building

we are going to train Linear Regression Model; We need to split out data into two variables x and y where x is independent variable (input) and y is dependent on x (output) we could ignore address column as it is required for our model

```
In [15]: x=df1[['cgpa', 'placement_exam_marks' ]]  
y=df1[ 'placed']
```

To Split my dataset into training and test data

```
In [16]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [17]: from sklearn.linear_model import LinearRegression  
lr= LinearRegression()  
lr.fit(x_train,y_train)
```

Out[17]: LinearRegression()

```
In [18]: lr.intercept_
```

Out[18]: 0.3934223238880175


```
In [19]: coeff = pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

Out[19]:

	Co-efficient
cgpa	0.011230
placement_exam_marks	0.000127

```
In [20]: prediction = lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

Out[20]: <matplotlib.collections.PathCollection at 0x1a384bfd460>



```
In [21]: lr.score(x_test,y_test)
```

Out[21]: -0.007364949116704933